

Applying fault tolerant Safra to Chandy Misra

Technical Report

Per Fuchs

August 7, 2018

1 Introduction

This technical report presents an experiment using our fault tolerant Safra version (short SafraFT) (**ourpaper**) to detect termination of a fault tolerant Chandy Misra routing algorithm. The fault-tolerant Chandy Misra version is developed for this project as an extension of the fault sensitive Chandy Misra algorithm described in (Fokink, 2018) on page 57. An explanation of the extension and further important implementation decisions can be found in section 2.

The experiment aims to

- backup our claim that SafraFT is correct by using it in a realistic setting and verifying that termination is detected in a timely manner after actual termination occurred
- compare the performance of SafraFT to a fault sensitive Safra implementation (abbreviated SafraFS) from (Demirbas & Arora, 2000)
- demonstrate the ability of SafraFT to handle networks of 25 and up to 2000 nodes in a fault-free, 1 to 5 faults and highly faulty (90% node failure) environment

Towards this aim, I measure the following dependent variables

- total number of tokens send
- number of tokens send after the basic algorithm terminated
- number of backup tokens send
- average token size in bytes
- processing time for Safra’s procedures
- time spent processing the basic algorithm’s procedures
- wall time for the complete computation
- wall time between termination of the basic algorithm and detection of the fact

All these metrics are measured within the following environments:

- network size of 25, 250, 500, 1000 and 2000 nodes
- using SafraFS and SafraFT
- in a fault-free environment
- for SafraFT additionally with 1 - 5 and up 90% node failures (simulating nearly fault-free networks and highly faulty environment)

I do not aim to show the exact relationships between the dependent variables and the independent variables e.g. the relationship between backup tokens send and the number of faults. This is because the exact relationship depends heavily on the basic algorithm, the network and even the hardware the system is running on. Therefore, detailing the dependence would not be helpful to anyone considering to apply our algorithm to his system. However, this experiment should enable the reader to judge if SafraFT could be used for his system and convince him that SafraFT performance is comparable to that of SafraFS in a fault-free environment (except for its higher bit complexity). Furthermore, this report aims to show how SafraFT behaves in a faulty environment.

Before performing this experiment George applied SafraFT to a simulated basic algorithm in a multi-threaded environment emulating a distributed system. These experiments showed strong evidence towards correctness and scalability of SafraFT. The implementation, technical report and results can be found here: (**georgework**).

The experiment presented here is performed on a reviewers recommendation to add ‘compelling end-to-end application’.

2 Methods

This section describes the most important software design decision of this experiment e.g. how a fault tolerant Chandy Misra algorithm can be implemented or how I simulate faults. Later, I describe the software and machines used.

2.1 Chandy Misra

- src Distributed algorithms by Fokking - either write and connect this or just reference this in the next chapter

2.2 Fault tolerant Chandy Misra

The fault-tolerant Chandy Misra version used for our experiments constructs a sink tree in an undirected network under the assumption that a perfect failure detector is present at each node (a detector that does not suspect nodes that haven't actually failed and detects each failure eventually). Other than the original Chandy Misra algorithm, the fault tolerant version requires FIFO-channels. Furthermore, I assume that the root node cannot fail because otherwise there is no sink tree to construct.

As far as Chandy-Misra is concerned nodes are only interested in crashes of their parents and other ancestors on their path to the root node. If a node **X** detects a crash of its parent, it sends a **REQUEST** message to each neighbour. If a neighbour **Y** of **X** receives a **REQUEST** message, it answers with a **DIST d** message where **d** is its own distance. To save messages **DIST d** is only sent if $d < \infty$. If **Y** happens to be a child of **X**, it resets its own **dist** and **parent** value to ∞ respectively \perp and sends a **REQUEST** message to all its neighbours. In this case **Y** sends no **DIST** message as answer to **X**.

The requirement for FIFO channels is best understood by a counterexample based on a non-FIFO network. The network used for this example is shown in fig. 1. Only important messages are mentioned; all others can be assumed to be sent and received in any order. The Chandy Misra algorithm starts with **A** as root node sending **DIST 0** messages to **B** and **C** which on receive consider **A** their parent and update their **dist** variable. They also send **DIST** messages to their neighbours. When **C** and **D** receive the **DIST** messages from **B** respectively **C** they consider **B** respectively **C** their parents. **C** sends a **DIST 2** message to **D** - let's call it **M1**. Now **B** crashes and when **C** detects this, it sends a **REQUEST** message towards **A** and **D** - call the latter **M2**. If **M2** overtakes **M1**, **D** resets its variables and on receive of **M1** considers **C** its parent which is correct but with an incorrect **dist** value of 2. All **DIST** messages received by **D** from now on have a higher distance value and are dismissed. So the error is never corrected.

A straightforward fix for this is to use FIFO networks because the original problem is that **M2** overtook **M1**.

SafraFT uses the same FIFO channels as its basic algorithm during the experiment. Nonetheless, we still claim it does not need this property. This claim is straightforward to prove: SafraFT guarantees that at most one message is in any channel at all times because it forwards a single token and when a backup token is issued, it is sent to a different node than the original token and only the first of both tokens are forwarded afterwards. Following the reasoning that only one message is in flight between any two nodes, SafraFT is indifferent to the property FIFO property of the channels.

I argue that this augmented Chandy Misra algorithm constructs correct sink trees in presence of fail-safe failures. Each failure only affects nodes that see themselves as children, grandchildren or deeper ancestor of the crashing node; that is, a failure only affects subtrees. The children eventually send **REQUEST** messages to all their neighbours because the perfect failure detector guarantees that each node failure is eventually detected by them. The neighbours send **REQUEST** messages to all their neighbours if they receive a **REQUEST** message from their parent. Therefore, eventually, all nodes in the subtrees of a failing node are reached by a **REQUEST** message and reset their **dist** and **parent** values. Also, all neighbours that receive a **REQUEST** message from a node that is not their parent, answer it with their current **dist** value. This allows nodes in the affected subtree to rebuild




Figure 1: None FIFO network with A as root to demonstrate necessity of FIFO networks.

new paths toward the root node. These new paths are correct when the answering node is not part of any affected subtree. However, if they are part of an affected subtree (e.g. grandchildren of the crashed nodes), invalid paths are introduced - as these nodes might have not been reached by any **REQUEST** message and therefore still believe they have a valid path towards the root. These invalid paths are corrected when the grandchildren are reached by the **REQUEST** message of their parent because on the receipt they send **REQUEST** messages which reset all nodes considering them parents. This behaviour of introducing invalid paths that are corrected later might lead to a bad theoretical message complexity but did not hinder the experiments.

The presented fault tolerant Chandy Misra algorithm could be incorrect as it calculated the wrong tree in Unfortunately, I could not fix it in time for the final experiment runs. However, this does not influence the evaluation of SafraFT because the correctness check for SafraFT and Chandy Misra in the experiment setup are designed to work independently from each other (see ??). The verification of SafraFT does show that SafraFT detected termination correctly in these cases and did not cause the corrupted Chandy Misra result.

Additionally, to fixing the last bugs, fault-tolerant Chandy Misra version could be improved by relieving the necessity for FIFO-channels, a more formal and adapted proof of correctness and a thorough complexity analysis.

2.3 Fault Simulation

I simulate faults by stopping Safra and the basic algorithm on the failing instance. In particular, a crashed node does not send or react to tokens or basic messages. Faults are triggered before and after interesting events e.g. directly before or after sending a basic message or token. Before every experiment run it is determined at random which node is going to fail, on which event it is going to fail and after how many repetitions of this event e.g. after the 3rd time it forwards a token.

In particular, I selected the following events to trigger a crash if not specified differently the crash is triggered before or after the event:

- sending a token (1 to 3 repetitions)
- sending a backup token (1 to 2 repetitions)
- before receiving a token (1 to 3 repetitions)
- sending a basic message (1 to 5 repetitions)

The range of repetitions is limited to maximize the chance that a node meant to fail actually does so. However, a node that is planned to fail is not guaranteed to do so. For example, this leads to runs where 90% of the nodes should fail but only 88% do so. I verify for every run that the number of crashed nodes does not vary more than 10% from the expected number. For runs with only 1 to 5 failing nodes, I confirm that at least one instance failed.

Alternatively, to the chosen approach to trigger failures, I considered the more random mechanism of running a thread that kills an instance after a random amount of time. One could argue that this would be more realistic. However, I believe this kind of an approach leads to less interesting failures because the vast majority of these failures would occur during idle time. Furthermore, most failures between internal events are observed as exactly the same on all other nodes. For example, other nodes cannot observe if a failure happened before an internal variable changed or after. In fact, they can only observe a difference when the failure happens after or before sending a message to them. Hence, I have chosen a fault mechanism that focusses on these distinguishable scenarios. As one might notice, the failure points are chosen to give rise to many different situations for our Safra version to deal with. I deliberately decided against choosing special failure points with regard to the basic algorithm because this would lead to less focused testing of the fault tolerance of Safra.

2.4 Fault Detection

Our Safra version assumes the presence of a perfect fault detector. This kind of fault detection is easy to implement and integrate with the system e.g. (Fokkink, 2018) on page 113 describes a straightforward implementation.

As building a perfect fault detector is a well known and solved problem, but nonetheless, time-consuming, I decided to avoid implementing one. For this experiment, fault detection is simulated by sending CRASH messages from crashing nodes to their neighbours. These crash messages are sent through different channels than basic and Safra messages because otherwise they would arrive in FIFO order with them and this would exclude situations where a basic message is received after the crash of the sender has been detected.

CRASH messages are not broadcasted to all nodes because IBIS (the message passing library I used) does not provide broadcasting.

2.5 Offline Analysis

For the experiment, I measure some metrics before and after termination e.g. the total token count and tokens sent after termination of the basic algorithm. To allow generation of these metrics, I need a close estimate of when the basic algorithm terminated. For this, I generate log files of events during execution and analyse these afterwards.

Termination is commonly defined by:

1. All nodes are passive
2. No messages are in the channels

To allow verification of 1. every node logs changes of its active status. The second point can be verified indirectly by logging all changes of the message counters managed by Safra’s algorithm. These counters are incremented for each message sent and decremented when a message is received. Therefore, one can conclude that no messages are in flight when the sum of all counters is 0. All nodes log the aforementioned events combined with a timestamp. By sorting and keeping track of active statuses, as well as, the sum of message counters, one can estimate the time of basic termination by the timestamp of the last node becoming passive while the sum of all message counters is zero. This technique is similar to the one Safra uses but a global view on the system achieved by offline analysis allows to detect the time of basic termination more closely.

With this system in place, it is possible to determine the number of tokens sent after termination by logging each token sent event and categorizing them during the offline analysis.

Processing time metrics are determined by the same principle: processing time is logged online and is grouped into total and after termination by analysing the logs after the run. Wall time between basic termination and detection by Safra is determined by comparing the timestamp of the event causing termination with the timestamp of the to announce call by Safra.

2.6 Environment

This chapter describes software, hardware and simulated network topology used for the experiments.

2.6.1 IBIS

IBIS is a Java-based platform for distributed computing developed by researchers of the Vrije Universiteit Amsterdam. Its subpart IPL is used as the message passing library for this project. I use version 2.3.1 which can be found on GitHub: <https://github.com/JungleComputing/ipl/releases/tag/v2.3.1>.

Communication channels in IPL are backed by TCP and provide asynchronous messaging. For this experiments, I also used IPL's ability to guarantee FIFO channels.

2.6.2 DAS 4

The experiment is conducted on the part of DAS-4 that belongs to the Vrije Universiteit Amsterdam. The nodes use primarily SuperMicro 2U-twins with Intel E5620 CPUs and a Linux CentOS build with the kernel version 3.10.0-693.17.1.el7.x86_64. For communication 1Gbit/s Lan is used. At the time of the experiments, the VU had 41 nodes with 8 cores each. Therefore, multiple instances were run on each physical node to be able to test our algorithm on decent sized networks (the number of machines and instances per machine can be found in table 1). This is possible because Safra and Chandy Misra are both communication heavy with rather low processing and memory requirements.

2.6.3 Network Topology

Chandy Misra needs a network topology to work on. It requires an undirected network. To generate more interesting runs I use weighted networks.

Our Safra version needs an undirected ring. For simplicity in my setup, this ring is part of the network the basic algorithm runs on. That means there is always an undirected ring connecting all nodes within simulated networks.

All networks for the experiments are generated by choosing randomly between 1 and 5 neighbours for each node and assigning a random, unique weight between 1 and 50000 to each channel.

After this channels with the heavy weight of 400000 are added between the root and some nodes to ensure the network stays connected when nodes fail. For this, I calculate the expected network with knowledge of the nodes predetermined to fail and add connections between nodes that could become disconnected as some other nodes fail and the root. These channels are heavyweight to avoid using them over 'regular' channels which might create highly similar topologies with a lot of nodes directly connected to the root.

At last, channels are added to form an undirected ring in which each channel has the weight of 100000. Again the weight is chosen to avoid 'overusing' the ring channels for the trees built.

For the experiments, I am not interested in the relationship between network topology and the measured metrics because this kind of analysis seems too detailed to show the correctness of our Safra version or to compare our Safra version with the traditional Safra version. Therefore, the exact network topology is not recorded. Anyhow, I output the depth of resulting trees and the nodes per level for runs up to 500 nodes. This information is recorded to allow for quality checks and ensure the experiments were not run on trivial networks. Generating these statistics for runs with over 500 nodes turned out to be time-consuming. Hence, I decided not to generate them. However, if the network topology leads to none trivial, deep trees for 500 nodes and less, it seems highly unlikely that it would not for even more nodes. Furthermore, I verified that the topology generated for more than 500 nodes still leads to interesting cases by spot sample.

Algorithm	Network	Faults	Repetitions	#Instances / DAS-4 node
SafraFS	50	0	100	25
SafraFS	250/500/1000/2000	0	100	125
SafraFT	50	0	100	25
SafraFT	250/500/1000/2000	0	100	125
SafraFT	50	5n	100	25
SafraFT	250/500/1000/2000	5n	100	125
SafraFT	50	90%	100	25
SafraFT	250/500/1000/2000	90%	100	125

Table 1: List of all configurations run. Per physical DAS-4 node with 8 cores multiple virtual instances in their own processes where run (column ‘#Instances / DAS-4 node’). The amount of physical node equates to network size divided by instances.

3 Results

Over the next sections I present the main results of my experiment to support our claim that SafraFT is correct, to show how SafraFT compares to SafraFS and to exemplify the performance of SafraFT under the presence of faults.

The observations are based on runs of the system described in section 2 on the DAS-4 cluster at the Vrije Universiteit of Amsterdam. I measured runs on networks from 50 to 2000 nodes for SafraFT and SafraFS. SafraFT was also tested with 1 to 5 nodes failing per run (dubbed 5n) and with 90% node failure. table 1 presents how many repetition of each configuration were run.

The raw data including a manual how to interpret it can be found [TODO](#) here.

3.1 Correctness of SafraFT

The experiment is aimed to support our paper with a practical, correct application of our algorithm. Towards this goal I build multiple correctness checks into the experiment.

To assure nothing happens after termination detection, the application logs if any messages are received or actions are executed after termination has been detected and announced. The analysis tools provided with the experiment point these logs out to make sure they are not overlooked.

To proof termination is not detected to early, I use offline analysis (see section 2.5) to determine the point of actual termination and verify that detection happened after. All

However, the experiment revealed that the framework for termination chosen to develop SafraFT is not complete and does not cover all cases for my experiment setup. SafraFT is developed for the following and commonly used definition of termination:

1. All nodes are passive
2. No messages are in the channels

This definition is based on the fact that a node is either an initiator or can only become active if it receives a message. However, under the presence of failures and if additionally the outcome of the algorithm depends on the set of alive nodes, nodes might get activated by the detection of a failure. For example, when Chandy Misra builds a sink tree, nodes that detect a crash of their parents will become active afterwards to find a new path towards root. The idea described above leads to the following concrete scenario: lets consider the situation that all node are passive and no message are in the channel. In other words, the system terminated by our definition. Node X forwards the token and crashes afterwards. Node Y calls announce after receipt of the token. Assume node Z is a child of X and detects the crash of its parent, it becomes active after termination has been formally reached and announced. By sending out REQUEST message, it might activate other nodes again.

To conclude, the definition of termination that our algorithm is build upon does not fully capture our choice of basic algorithm which could lead to an early detection of termination.

To verify if situations like this actual occur during the experiment, I analysed the logs generated according to both definitions of termination - the one used to develop SafraFT and the same definition but extended with the assumption that termination is postponed until the last crash leading to activity in the basic algorithm has been detected.

As stated above, SafraFT did never detect termination too early according to the first definition. According to the extended definition it detects termination too early in. However, due to fact that repairing the sink tree after detecting a parent crash is quite fast and there is a short time window to do so while the announce call propagates to all nodes only in termination.

I carefully reviewed each repetition in which termination is detected too early according to the extended definition to verify that early termination detection is in fact caused by a situation as described above. The logs of these runs provide a summary of all detections of parent crashes close to the announce call to ease this procedure.

3.2 Comparison of Safra versions

This section compares SafraFS and SafraFT. Additionally, it analysis how the network size influences both algorithms.

The number of tokens send in total and after termination is presented in fig. 2.

The key observation is that SafraFS and SafraFT behave highly similar. There is no notable difference between the medians, the variance and the ratio between total tokens and tokens after termination for the same network size - except for a unusual high variance in tokens after termination SafraFS for 1000 nodes and a high variance in total tokens for SafraFS for 2000 nodes. As one would expect, the number of tokens seems to grow linearly with the network size. Note that the first network size is 5 times smaller than the second for bigger networks the size doubles for each run.

The bit complexity of SafraFS is constant. In this experiment each token of SafraFS contains 12 bytes. SafraFT has a bit complexity linearly to the network size (when no faults occur). For a network of 50 nodes each token has 420 bytes; a token in a 2000 node network counts 16020 bytes. The growth can be described by $bytes = 8 * < networksize > + 20$.

I measured two kinds of timing metrics in this experiment. On the one hand, there are the wall time metrics of total time and total time after termination. Both were recorded in elapsed seconds between two events. These events are start of the Safra and basic algorithm until each instance is informed of termination for total time. Total time after termination is defined as the ammount of seconds between the actual termination (extended termination definition from section 3.1) and the event of an node calling announce. On the other hand, there are basic, Safra and Safra after termination processing times (including the time needed to send of messages). These are the accumulated times all instances needed to process basic or Safra functionality. Total times and processing times are measured in a different way and should not be compared directly for multiple reasons. First, while total times include idle times, time spent for logging (where the process did not execute or methods), processing time do not include these. Secondly, total time is wall time between two events and processing times are accumulated over all processes. One particular example for when this leads to differences is that time spent concurrently by two processes counts double in processing time metrics but only once in wall time metrics.

One can observe in table 2 that SafraFT uses much more processing time than SafraFS and most of this time is spent between actual termination and termination detection. Furthermore, one sees that while SafraFS timing grow much slower in bigger networks than SafraFT timings and that the difference becomes bigger. This hints for a change in time complexity between the two versions.

A small subexperiment of excluding the time spent to send messages from the processing time revealed that most of this difference can be tributed to writting tokens onto the wire. As we know from a previous paragraph, the number of token send does not differ between the algorithms. Therefore, I believe that these differences are caused by the higher bit complexity of SafraFT. This would explain the total increase in the timing from SafraFS to SafraFT, as well as, the change of time complexity. The time complexity would change because the increase in network size leads to more token being sent (as in SafraFS) but also to bigger tokens being sent.

Network	Basic	Safra FS	Overhead FS	Safra FT	Overhead FT
50	1.248	0.035 (0.014)	2.8	0.059 (0.026)	4.73
250	11.202	0.200 (0.069)	1.79	0.469 (0.228)	4.19
500	62.525	0.488 (0.128)	0.78	1.304 (0.740)	2.09
1000	185.131	1.057 (0.228)	0.57	3.812 (2.486)	2.06
2000	1190.031	3.100 (0.465)	0.26	13.794 (9.005)	1.16

Table 2: Total processing times and processing times after termination (in braces) in seconds and overhead caused by processing Safra’s methods in percent

Network	Safra FS	Safra FT	Δ	Safra FS after	Safra FT after	Δ
50	0.219	0.232	1.06	0.048	0.077	1.6
250	1.011	1.333	1.32	0.35	0.665	1.9
500	2.168	3.178	1.47	0.726	1.699	2.34
1000	3.776	7.124	1.89	1.302	4.634	3.56
2000	9.711	19.593	2.02	2.664	11.666	4.38

Table 3: Wall times total and after termination for SafraFT and SafraFS in seconds with ratios

Most of the additional processing time for SafraFT is spent after termination. By the hypothesis introduced in the last paragraph, this is explained by most tokens being sent after termination (see fig. 2).

The processing time table 2 also presents a comparison of the time spent for the basic algorithm and both Safra versions. Although, SafraFT uses significantly more time, the overhead on the processing time stays moderate with a maximum of

The same pattern of SafraFT using more time and reacting stronger to an increase in network size is visible for total times in table 3. Again the majority of this time is spent after actual termination for the same reason as before. I would like to note that the low processing time overhead of Safra (maximum of is not in contradiction to the large amount of wall time spent after termination (maximum of These seemingly opposing results arise from the difference between wall time and processing time: the basic algorithm is much more active in the beginning that is when it accumulates a lot of processing time; while Safra causes a lot of idle time at the end when all processes wait for their predecessor to pass on the token. This idle time is not included in processing time but wall time does include it.

To conclude, the experiments confirm that the message complexity of SafraFT remains as for the fault sensitive version but its higher bit complexity causes a higher time complexity which leads to a later termination detection. Still, SafraFT causes only a moderate processing time overhead between

3.3 Influence of faults

In the following paragraphs, I present and explain the data generated by runs under the presence of node crashes. I run to really different scenarios: one considering networks with 1 to 5 nodes failing and one with 90% of all nodes crashing. These scenarios are chosen to show SafraFT in both the realistic case of a low number of faults to handle, as well as, the an extreme case; with the aim to confirm that SafraFT handles both cases correctly and without unreasonable deterioration in any metric.

Network	No faults	5n	Δ	90%	Δ
50	61 (44)	102 (54)	1.67 (1.23)	106 (13)	1.74 (0.30)
250	268 (241)	433 (305)	1.62 (1.27)	544 (64)	2.03 (0.27)
500	529 (488)	976 (661)	1.84 (1.35)	1086 (124)	2.05 (0.25)
1000	1032 (985)	1586 (1308)	1.54 (1.33)	2174 (225)	2.11 (0.23)

Table 4: Tokens before and after termination (in braces). For different fault scenarios compared to fault-free networks.

3.3.1 Tokens

For both the networks where between 1 and 5 nodes failed, as well as, for the highly faulty runs with 90% node failure, the number of tokens increased compared to runs without any faults. More so highly faulty networks more faults. The data is presented fig. 3 and table 4.

Otherwise, the two types of fault simulation had a highly different influence on the tokens and tokens after termination metrics.

Between 1 and 5 node failures lead to a strong increase in the variance of tokens sent, as well as, tokens sent after termination. This seems reasonable as runs with failing node might lead to more different situations as runs without fails e.g. one failing node could easily cause an extra token when it leads to an backup token being issued and forwarded (this token is marked black until for a whole run), at the same time, a single failing node that is a leaf in a Chandy Misra sink tree and that crashes just before the call of announce at the successor causes no further activity. The same example provides an idea why the variability increases in big networks. That is because one extra round in a big network has a much higher impact on the token count than in a small network.

A example crafted similarly would explain the extremely high maximum of tokens sent and overlinear increase of the average number of tokens in networks with 2000 nodes for 1 to 5 nodes failing.

Other than networks with one to 5 failing nodes, networks with up to 90% node failure lead to a similar variance in tokens and tokens after termination as a fault free networks. A likely explanation is that the low survival rate of 1 out of 10 leads to less different scenarios than 995 nodes surviving as in the fault scenario treated the last paragraphs.

Even though, only one 10th of the nodes survive to participate in the latter token rounds, the highly faulty networks produced more tokens than any other network of the same size. That is most likely due to the high amount of backup tokens generated (also shown in fig. 3)

Different from all other networks, highly faulty networks exhibit a much lower token to token after termination ratio caused by the low number of nodes alive in the last rounds.

3.3.2 Backup tokens

The average amount of backup tokens sent for either fault simulation and every network size is lower than the number of faults. This is because SafraFT only issues backup tokens when the fault of its successor is detected via the fault detector but not if this fault is noticed by receiving a token. That result is surprising because the simulated fault detector in my setup detects faults in a really timely manner. Still, there are some runs where 1 to 5 nodes fail but no backup token is issued for all network sizes.

The other extreme where more backup tokens are issued than faults occur exists as well. This can be explained by my decision to have node failing after issuing a backup token. For example, nodes A, B and C follow each other in the ring, node C fails which is detected by B and a backup token is issued. After, issuing the backup token B fails on detection A issues a backup token towards C. Only then A detects the failing of C and issues a second backup token to its new successor.

Network	No faults	5n	Δ	90%	Δ
50	0.059 (0.026)	0.085 (0.031)	1.44 (1.19)	0.159 (0.013)	2.69 (0.50)
250	0.469 (0.228)	0.643 (0.274)	1.37 (1.20)	1.709 (0.066)	3.64 (0.29)
500	1.304 (0.740)	1.956 (0.967)	1.50 (1.31)	3.972 (0.165)	3.05 (0.22)
1000	3.812 (2.486)	5.296 (3.258)	1.39 (1.31)	8.333 (0.386)	2.19 (0.16)

Table 5: Processing times before and after termination (in braces) in seconds. For different fault scenarios compared to fault-free networks.

Network	No faults	5n	Δ	90%	Δ
50	0.232 (0.077)	0.289 (0.072)	1.250 (0.940)	0.418 (0.023)	1.800 (0.300)
250	1.333 (0.665)	2.163 (0.680)	1.620 (1.020)	3.537 (0.122)	2.650 (0.180)
500	3.178 (1.699)	4.343 (1.848)	1.370 (1.090)	6.834 (0.280)	2.150 (0.160)
1000	7.124 (4.634)	9.308 (5.739)	1.310 (1.240)	12.696 (0.678)	1.780 (0.150)

Table 6: Total times before and after termination (in braces) in seconds. For different fault scenarios compared to fault-free networks.

3.3.3 Processing Time

The observations of this chapter are backed by table 5.

As for tokens, one can see an increase in processing times before termination under the presents of faults compared to fault free runs. This increase is stronger for runs with more faults. Which is no surprise as these runs also produced more tokens.

For runs with 1 to 5 failing node a higher variability in the processing times before and after termination becomes apparent. A likely reasons for this has been explained in the section 3.3.1. As the processing time grows for this runs, so does the processing time after termination.

For highly faulty networks one observes results in line with the results for tokens in these networks. The variability for processing time before or after termination is not raised compared to fault free networks. The processing time taken is even higher than the one for less faulty networks.

Less processing time after termination is spent than for fault free networks because less tokens need to be send.

3.4 Total time

In line with the observations from the two sections before, one observes:

- an increase in total time spent which is more pronounced on more faulty networks
- a higher variability for time spent before and after termination for less faulty networks
- less time spent after termination by highly faulty networks

The absolute numbers and relative increases are presented in table 6.




Figure 2: Tokens (blue) and tokens after termination (orange) for SafraFT and SafraFS in fault-free runs for all network sizes.




Figure 3: Tokens (blue) and tokens after termination (orange) for SafraFT on the x-axis for 5 network sizes and 5n and 90% fault configurations on the y-axis. Backup tokens in green, shown only for 90%

4 Discussion

why does this agree with our hypotheses or not? point out and explain anomalies

Limitations of my approach improvements of my approach

What do my results imply about Safra?

Comparison to George's experiments?

Our setup of a basic algorithm completing its work relatively fast and termination detection taking some time afterwards clearly shows a drawback of fault tolerant Safra. However, a system with a long running basic algorithm e.g. multiple hours, would put this time in a whole different perspective. Then the seconds taken to detect termination would be less of an issue and the moderate processing time overhead demonstrated far more important.

References

- Demirbas, M. & Arora, A. (2000). *An optimal termination detection algorithm for rings*. Technical Report OSU-CISRC-2/00-TR05, The Ohio State University.
- Fokkink, W. (2018). *Distributed algorithms: an intuitive approach*. The MIT Press.